

# MOHIT GUPTA

Mathura, Uttar Pradesh, India – 281003

Mobile: +91 7300554546 | [mohit.gupta.cr1@gmail.com](mailto:mohit.gupta.cr1@gmail.com) | [LinkedIn](#)

## Summary

Data Scientist with 2.5 years of experience in NLP, LLM fine-tuning, and AI agents. Skilled in Python, Docker, and microservices, with a focus on AI infrastructure, automation, and aligning solutions with business goals.

## Technical Skills

- **Programming:** Python (PyTorch, TensorFlow)
- **ML & AI:** LLMs, Transformers, NLP, Multi-Agent Systems, RAG, LangChain, LangGraph
- **Infrastructure:** Docker, Flask, FastAPI, Kubernetes, Elasticsearch, Nginx
- **Cloud:** GCP, Microsoft Azure
- **Databases:** MongoDB, SQL
- **Tools:** Pandas, NumPy, Scikit-learn, Selenium, BeautifulSoup

## Experience

### Data Scientist

IntelleWings Pvt. Ltd. | Chandigarh | Mar 2022 – Sep 2024

- Met data requirements in alignment with business needs provided by clients, ensuring accuracy and efficiency in deliverables.
- Developed NLP pipelines processing 500K+ documents daily, reducing processing time by 60%.
- Engineered scraping infrastructure to extract and process over 500k+ data points daily with high reliability and uptime.

## Projects

### Multi-Agent News Intelligence System *Python, FastAPI, Docker, GCP, FAISS, Sentence Transformers, Gemini API*

[GitHub](#) | [Live Demo](#)

- Production dashboard analyzing 600+ news articles daily with semantic search & automated GCP deployment.
- Designed modular multi-agent pipeline (SearchAgent, ReportingAgent, Orchestrator) for AI-cited reports.
- Built async FastAPI services + FAISS retrieval for real-time queries with <2s response time.

### LLM-based Information Extraction *Python, PyTorch, Mistral 7B, BitsAndBytes, QLoRA, RAG, Transformers*

- Used Mistral 7B fine-tuned with QLoRA & BitsAndBytes for structured extraction from financial PDFs.
- Integrated RAG for context-aware responses, improving relevance and reducing manual review.
- Optimized prompt strategies for accuracy and latency, achieving 95% precision on financial data fields.

### Financial Risk API Development *Python, Flask, MongoDB, Elasticsearch, Docker*

- Built & maintained 100+ Flask APIs for real-time risk data with 24/7 uptime.
- Hybrid storage: MongoDB for docs, Elasticsearch for retrieval at scale.
- Developed crawlers for continuous data updates and monitoring for API performance.

## Education

### B.Tech, Computer Science

GLA University | 2019–2023